

The top banner features the Fudan University logo on the left, which includes the university's name in English ('FUDAN UNIVERSITY') and Chinese ('復旦大學'), along with the founding year '1905'. To the right of the logo are several interlocking blue gears. Some of these gears contain white icons: a building, a graduation cap, a medical cross, a person silhouette, and an atomic symbol. The background of the banner is light blue with a dark blue curved shape on the far right.

# Big Data Analytics & Applications


Bin Li

School of Computer Science


Fudan University

# ML Problems in BDA


- In BDA, there are two commonly seen ML problem settings
  - ▣ Supervised learning: **Classification**, **Regression**
  - ▣ Unsupervised learning: **Clustering**, **Dimensionality Reduction**



Supervised Learning  
Algorithms



Unsupervised Learning  
Algorithms



Semi-supervised  
Learning Algorithms

# Classification Problem

## ■ Inputs:

- ▣ Instances:  $x_1, x_2, \dots \in X$ , where  $X$  is feature space
- ▣ Class labels:  $y_1, y_2, \dots \in Y$ , where  $Y = \{1, 2, \dots, L\}$  is label set

## ■ Classification problem setting:

- ▣ Training data:  $\{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$
- ▣ Test data:  $\{(x', ?), (x'', ?), \dots, \}$

- **Objective:** Learn a function  $f: X \rightarrow Y$  on the training data such that  $f(x_n) = y_n$  for  $n = 1, 2, \dots, N$

$$\min_f \frac{1}{N} \sum_{n=1}^N \textit{classification\_loss}(f(x_n), y_n)$$

# Classification Problem

- Supervised learning example: Image classification


- ▣ Instances:  $x_1, x_2, \dots$  are images

- ▣ Class labels:  $y_1, y_2, \dots \in \{dog, cat, rabbit\}$

$f(\text{dog emoji}) = "dog"$

$f(\text{cat emoji}) = "cat"$

$f(\text{rabbit emoji}) = "rabbit"$



# Classification Problem

- Application provides the dataset  $\{(x_1, y_1), \dots, (x_N, y_N)\}$ 
  - Image annotation:  $x_n$  pixel image
  - Document categorization:  $x_n$  bag-of-words representation
  - User classification:  $x_n$  user profile
- $f: X \rightarrow Y$  is a selected model for certain applications
  - Nearest Neighbors Classifier
  - Linear Classifier
  - Logistic Regression
  - Support Vector Machine
  - Multilayer Perceptron
  - Decision Tree

# Regression Problem

## ■ Inputs:

- ▣ Instances:  $x_1, x_2, \dots \in X$ , where  $X$  is feature space
- ▣ Targets:  $y_1, y_2, \dots \in Y$ , where  $Y \subset \mathbb{R}^m$  is target domain

## ■ Regression problem setting:

- ▣ Training data:  $\{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$
- ▣ Test data:  $\{(x', ?), (x'', ?), \dots, \}$

- **Objective:** Learn a function  $f: X \rightarrow Y$  on the training data such that  $f(x_n) = y_n$  for  $n = 1, 2, \dots, N$

$$\min_f \frac{1}{N} \sum_{n=1}^N \text{regression\_loss}(f(x_n), y_n)$$

# Regression Problem

- Supervised learning example: Income Prediction


- ▣ Instances:  $x_1, x_2, \dots \in \text{Education} \times \text{Seniority}$

- ▣ Targets:  $y_1, y_2, \dots \in \text{Income}$

$$f(12, 3) = 50K$$

$$f(16, 5) = 80K$$

$$f(19, 8) = 200K$$




# Regression Problem

- Application provides the dataset  $\{(x_1, y_1), \dots, (x_N, y_N)\}$ 
  - Stock trend prediction:  $x_n$  variates,  $y_n$  index
  - Location prediction:  $x_n$  previous locations,  $y_n$  new location
  - Rating prediction:  $x_n$  user profile,  $y_n$  rating
- $f: X \rightarrow Y$  is a selected model for certain applications
  - Nearest Neighbors Regression
  - Linear Regression
  - Support Vector Regression
  - Gaussian Process Regression
  - Multi-Layer Perceptron
  - Decision Tree




# Classification vs Regression

- It seems that classification and regression are similar
- What makes them different? - **Loss** ✕




$$\min_f \frac{1}{N} \sum_{n=1}^N \text{classification\_loss}(f(x_n), y_n)$$

$$\min_f \frac{1}{N} \sum_{n=1}^N \text{regression\_loss}(f(x_n), y_n)$$



# Classification Loss


- Zero-One Loss:  $l_{0-1}(f(x_n), y_n) = 1(y_n f(x_n) \leq 0)$
- Hinge Loss:  $l_{\text{hinge}}(f(x_n), y_n) = \max(1 - y_n f(x_n), 0)$
- Logistic Loss:  $l_{\text{logistic}}(f(x_n), y_n) = \ln(1 + \exp(-y_n f(x_n)))$



$$\min_f \frac{1}{N} \sum_{n=1}^N \text{classification\_loss}(f(x_n), y_n)$$

# Regression Loss

- Quadratic Loss (L2-Loss):  $l_{L2}(f(x_n), y_n) = (y_n - f(x_n))^2$
- Absolute Loss (L1-Loss):  $l_{L1}(f(x_n), y_n) = |y_n - f(x_n)|$
- Huber Loss:  $l_{Huber}(f(x_n), y_n) = \begin{cases} (y_n - f(x_n))^2, & |y_n - f(x_n)| \leq \delta \\ |y_n - f(x_n)|, & |y_n - f(x_n)| > \delta \end{cases}$




$$\min_f \frac{1}{N} \sum_{n=1}^N \text{regression\_loss}(f(x_n), y_n)$$



# Generalization

- Generalization error is a measure of how accurately a model is able to predict outcome values for previously **unseen** data ✕

Regression:






predictor too inflexible:  
cannot capture pattern



predictor too flexible:  
fits noise in the data

Classification:




# Regularization

- In ML, regularization is a process of introducing additional information in order to prevent **overfitting**
- Regularized Empirical Risk Minimization:

$$\min_f \frac{1}{N} \sum_{n=1}^N \text{loss}(f(x_n), y_n) + \lambda R(f)$$

- Some explanations for regularization
  - Solve ill-posed (**underdetermined**) problem
  - Impose **Occam's razor** on the solution
  - Impose certain **prior distributions** on model parameters

# Structural Risk Minimization



# Objective Function

- General form of objective function for supervised learning

$$\min_f \frac{1}{N} \sum_{n=1}^N \text{loss}(f(x_n), y_n) + \lambda R(f)$$

- Different combinations of  $f(\cdot)$ ,  $\text{loss}(\cdot)$ , and  $R(\cdot)$  will result different machine learning models

- Ordinary Linear Model: Ridge Regression

- Linear model:  $f(x_n) = w^\top x_n$
- Quadratic loss:  $\text{loss}(f(x_n), y_n) = (y_n - w^\top x_n)^2$
- L2-norm regularization:  $R(f) = ||w||^2$

$$\min_w \frac{1}{N} \sum_{n=1}^N (y_n - w^\top x_n)^2 + \lambda ||w||^2$$

# Ridge Regression

## ■ Ridge Regression

□ Linear model:  $f(x_n) = w^\top x_n$

□ Quadratic loss:  $loss(f(x_n), y_n) = (y_n - w^\top x_n)^2$

□ L2-norm regularization:  $R(f) = ||w||^2$

$$\min_w \frac{1}{N} \sum_{n=1}^N (y_n - w^\top x_n)^2 + \lambda ||w||^2$$

$$\Rightarrow \min_w (Y - Xw)^\top (Y - Xw) + \lambda w^\top w$$

← Convex Optimization  
(quadratic programming)

□ Let  $J(w) = (Y - Xw)^\top (Y - Xw) + \lambda w^\top w$  ✕

$$\frac{\partial J(w)}{\partial w} = -2X^\top (Y - Xw) + 2\lambda w = 0$$

$$w = (X^\top X + I\lambda)^{-1} X^\top Y$$



# Logistic Regression

## ■ Logistic Regression

□ Linear model:  $f(x_n) = w^\top x_n$

□ Logistic loss:  $loss(f(x_n), y_n) = -y_n \ln(\sigma_n) - (1 - y_n) \ln(1 - \sigma_n)$ ,  
where  $\sigma_n = \frac{1}{1 + \exp(-w^\top x_n)}$  and  $y_n \in \{0, 1\}$

□ L2-norm regularization:  $R(f) = ||w||^2$

$$\min_w \frac{1}{N} \sum_{n=1}^N -y_n \ln(\sigma_n) - (1 - y_n) \ln(1 - \sigma_n) + \lambda ||w||^2$$

□ Let  $J(w) = -\frac{1}{N} \sum_{n=1}^N (y_n \ln(\sigma_n) - (1 - y_n) \ln(1 - \sigma_n)) + \lambda ||w||^2$

$$\begin{aligned} \frac{\partial J(w)}{\partial w} &= \frac{1}{N} \sum_{n=1}^N -y_n(1 - \sigma_n)x_n + (1 - y_n)\sigma_n x_n + 2\lambda w \\ &= \frac{1}{N} \sum_{n=1}^N x_n(\sigma_n - y_n) + 2\lambda w \end{aligned}$$

# Maximum Likelihood Estimation

- In statistics, Maximum Likelihood Estimation (MLE) is a method of estimating the parameters of a statistical model, given observations  $D = \{(x_1, y_1), \dots, (x_N, y_N)\}$
- MLE attempts to find the parameter values that maximize the likelihood function ※

$$\max_{\theta \in \Theta} p(D|\theta) \Rightarrow \max_{\theta \in \Theta} \prod_{n=1}^N p(x_n, y_n|\theta) \Rightarrow \min_{\theta \in \Theta} \frac{1}{N} \sum_{n=1}^N -\log p(x_n, y_n|\theta)$$

- $p(x_n, y_n|\theta)$  is Gaussian distribution → Quadratic loss
- $p(x_n, y_n|\theta)$  is Laplacian distribution → Absolute loss
- $p(x_n, y_n|\theta)$  is Logistic function → Logistic loss

# Maximum *a posteriori*

- In Bayesian statistics, Maximum *a posteriori* (MAP) is an estimate of an unknown quantity, that equals the **mode** of the posterior distribution.
- MAP estimation can be seen as a regularization of MLE.


$$\arg \max_{\theta \in \Theta} p(\theta|D) = \arg \max_{\theta \in \Theta} \frac{p(D|\theta)p(\theta)}{\int_{\theta'} p(D|\theta')p(\theta')d\theta'} = \arg \max_{\theta \in \Theta} p(D|\theta) \mathbf{p}(\theta)$$

$$\begin{aligned} \max_{\theta \in \Theta} p(D|\theta)p(\theta) &\Rightarrow \max_{\theta \in \Theta} \prod_{n=1}^N p(x_n, y_n|\theta)p(\theta) \\ &\Rightarrow \min_{\theta \in \Theta} \frac{1}{N} \sum_{n=1}^N -\log p(x_n, y_n|\theta) - \mathbf{\log p(\theta)} \end{aligned}$$


- Compare to  $\min_f \frac{1}{N} \sum_{n=1}^N \text{loss}(f(x_n), y_n) + \lambda R(f) \times$

# ML Problems in BDA


- In BDA, there are two commonly seen ML problem settings
  - ▣ Supervised learning: **Classification**, **Regression**
  - ▣ Unsupervised learning: **Clustering**, **Dimensionality Reduction**



Supervised Learning  
Algorithms



Unsupervised Learning  
Algorithms



Semi-supervised  
Learning Algorithms

# Clustering Problem

- Instances:  $x_1, x_2, \dots \in X$ , where  $X$  is feature space
- **Objective:** Input instances  $x_1, x_2, \dots \in X$  and output corresponding cluster indicators  $z_1, z_2, \dots \in \{0,1\}^K$  for each instance, satisfying certain optimization criteria

$$\min_{\{z\}, \{\theta\}} \frac{1}{N} \sum_{n=1}^N \sum_{k=1}^K z_{n,k} \text{clustering\_loss}(x_n, \theta_k)$$

- $\theta_k$  denotes the parameter set of the  $k$ -th cluster
- $z_{n,k} \in \{0,1\}$  denotes whether the  $n$ -th instance belongs to the  $k$ -th cluster
- Goal: Find values of  $\theta_k$  and  $z_{n,k}$  to minimize the objective

# K-Means Clustering

- Minimize  $J = \frac{1}{N} \sum_{n=1}^N \sum_{k=1}^K z_{n,k} \|x_n - \mu_k\|^2$  in terms of the mean of the cluster  $\mu_k$  and the cluster indicator  $z_{n,k}$
- Given  $\{\theta_1, \dots, \theta_K\}$  fixed, since  $J$  is a linear function of  $z_{n,k}$ , this optimization can be easily obtained by



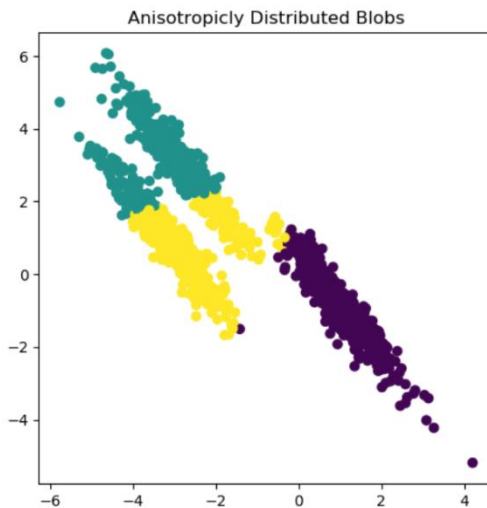
$$z_{n,k} = \begin{cases} 1, & \text{if } k = \arg \min_j \|x_n - \mu_j\|^2 \\ 0, & \text{otherwise} \end{cases}$$

- Given  $\{z_{1,1}, \dots, z_{N,K}\}$  fixed,  $J$  is a quadratic function of  $\mu_k$ , it can be minimized by setting its derivative w.r.t.  $\mu_k$  to zero

$$2 \sum_{n=1}^N z_{n,k} (x_n - \mu_k) = 0 \Rightarrow \mu_k = \frac{\sum_{n=1}^N z_{n,k} x_n}{\sum_{n=1}^N z_{n,k}}$$

# Assumptions of $K$ -Means

- Assumptions (sometimes limitations) in  $K$ -Means
  - ❑ Isotropically distributed
  - ❑ Distributed with equal variances
  - ❑ Clusters are evenly sized



# Gaussian Mixture Model (GMM)

- $K$ -Means is **hard-membership** clustering technique
  - $z_{n,k} \in \{0,1\}$ : Each instance can or cannot belong to a cluster
  - $\sum_{k=1}^K z_{n,k} = 1$ : Each instance can belong to only one cluster
- Is there a **soft-membership** clustering technique?
  - $z_{n,k} \in \{0,1\} \rightarrow \gamma(z_{n,k}) \in [0,1]$
  - $\sum_{k=1}^K \gamma(z_{n,k}) = 1$  still holds
- Gaussian Mixture Model

$$p(x_n|\theta) = \sum_{k=1}^K \pi_k \mathcal{N}(x_n|\mu_k, \Sigma_k)$$

- Anisotropically distributed
- Distributed with different variances
- Clusters are variously sized



# Maximum Likelihood of GMM

- Given  $\{x_1, \dots, x_N\}$  and we wish to model these instances using a GMM. The log likelihood function is

$$\ln p(X|\{\pi_k, \mu_k, \Sigma_k\}_k) = \sum_{n=1}^N \ln \left\{ \sum_{k=1}^K \pi_k \mathcal{N}(x_n | \mu_k, \Sigma_k) \right\}$$

- Set the derivatives of  $\ln p(X|\{\pi_k, \mu_k, \Sigma_k\}_k)$  w.r.t. the means  $\mu_k$  of the Gaussian components to zero

$$-\sum_{n=1}^N \frac{\pi_k \mathcal{N}(x_n | \mu_k, \Sigma_k)}{\sum_{k'=1}^K \pi_{k'} \mathcal{N}(x_n | \mu_{k'}, \Sigma_{k'})} \Sigma_k (x_n - \mu_k) = 0$$

$$\gamma(z_{n,k}) = p(z_{n,k} = 1 | x_n) = \frac{\pi_k \mathcal{N}(x_n | \mu_k, \Sigma_k)}{\sum_{k'=1}^K \pi_{k'} \mathcal{N}(x_n | \mu_{k'}, \Sigma_{k'})}$$

$$-\sum_{n=1}^N \gamma(z_{n,k}) \Sigma_k (x_n - \mu_k) = 0 \Rightarrow \mu_k = \frac{\sum_{n=1}^N \gamma(z_{n,k}) x_n}{\sum_{n=1}^N \gamma(z_{n,k})}$$

# Maximum Likelihood of GMM

- Set the derivatives of  $\ln p(X|\pi, \mu, \Sigma)$  w.r.t. the covariance matrix  $\Sigma_k$  of the Gaussian components to zero

$$\Sigma_k = \frac{\sum_{n=1}^N \gamma(z_{n,k})(x_n - \mu_k)(x_n - \mu_k)^\top}{\sum_{n=1}^N \gamma(z_{n,k})}$$

- Set the derivatives of  $\ln p(X|\pi, \mu, \Sigma)$  w.r.t. the mixing coefficients  $\pi_k$  subject to  $\sum_{k=1}^K \pi_k = 1$

$$\pi_k = \frac{\sum_{n=1}^N \gamma(z_{n,k})}{N}$$

- In the MLE of  $(\pi_k, \mu_k, \Sigma_k)$  we assume  $\gamma(z_{n,k})$  is given, which however is also conditioned on  $(\pi_k, \mu_k, \Sigma_k)$

# Expectation-Maximization (EM) Algorithm for GMM

- Initialize  $(\pi_k, \mu_k, \Sigma_k)$  and evaluate  $\ln p(X|\pi, \mu, \Sigma)$
- E-Step: Evaluate  $\gamma(z_{n,k})$  using the current  $(\pi_k, \mu_k, \Sigma_k)$

$$\gamma(z_{n,k}) = \frac{\pi_k \mathcal{N}(x_n | \mu_k, \Sigma_k)}{\sum_{k'=1}^K \pi_{k'} \mathcal{N}(x_n | \mu_{k'}, \Sigma_{k'})}$$


- M-Step. Estimate  $(\pi_k, \mu_k, \Sigma_k)$  using the current  $\gamma(z_{n,k})$

$$\begin{aligned}\pi_k &= \frac{\sum_{n=1}^N \gamma(z_{n,k})}{N} & \mu_k &= \frac{\sum_{n=1}^N \gamma(z_{n,k}) x_n}{\sum_{n=1}^N \gamma(z_{n,k})} \\ \Sigma_k &= \frac{\sum_{n=1}^N \gamma(z_{n,k}) (x_n - \mu_k)(x_n - \mu_k)^\top}{\sum_{n=1}^N \gamma(z_{n,k})}\end{aligned}$$

- Evaluate  $\ln p(X|\pi, \mu, \Sigma)$  and check the convergence

# EM Algorithm for GMM

## ■ Illustration of EM iterations for fitting a GMM



# EM Algorithm for Latent Variable Models

- Given a joint distribution  $p(X, Z|\theta)$  over observed variables  $X$  and latent variables  $Z$ , governed by parameters  $\theta$ , the goal is to maximize the likelihood function  $p(X|\theta)$  w.r.t.  $\theta$ .
- The general EM algorithm:
  - Initialize the parameters  $\theta^{\text{old}}$ ;
  - **E-Step**: Evaluate  $p(Z|X, \theta^{\text{old}})$ ;
  - **M-Step**: Evaluate  $\theta^{\text{new}}$  given by

$$\theta^{\text{new}} = \operatorname{argmax}_{\theta} \sum_Z p(Z|X, \theta^{\text{old}}) p(X, Z|\theta)$$

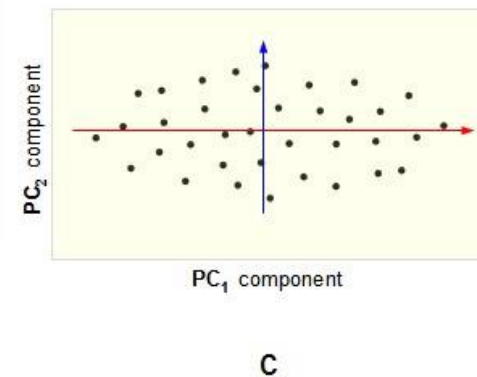


- Check the convergence of the parameter values; if convergence condition is not satisfied, set  $\theta^{\text{old}} = \theta^{\text{new}}$  and go to E-step.

# Latent Variable Models

- Latent Class Analysis (discrete latent variable model) is usually based on **cluster assumption**
  - Mixture Model
  - Topic Model
  - Relational Model
  - Latent Feature Model
  - etc.
- Latent Factor Analysis (continuous latent variable model) is usually based on **subspace assumption**
  - Principal Component Analysis
  - Matrix Factorization
  - etc.

# Principal Component Analysis

- Given  $\{x_1, \dots, x_N\} \in R^D$ , PCA is to project the data on to a space that maximizes the variance of the projected data
  - ▣ Noise filtering
  - ▣ Dimensionality reduction
  - ▣ Data visualization
  - ▣ Data compression
  - ▣ etc.



# Principal Component Analysis

- Suppose the first principal component is  $u_1$ , then each data point is projected onto a one-dimensional space  $u_1^\top x_n$
- The variance of the projected data is given by

$$\frac{1}{N} \sum_{n=1}^N (u_1^\top x_n - u_1^\top \bar{x})^2 = u_1^\top \left( \frac{1}{N} \sum_{n=1}^N (x_n - \bar{x})(x_n - \bar{x})^\top \right) u_1 = u_1^\top \Sigma u_1$$

- Maximize  $u_1^\top \Sigma u_1$  w.r.t.  $u_1$  subject to  $u_1^\top u_1 = 1$

$$\max_{u_1} u_1^\top \Sigma u_1 + \lambda_1 (1 - u_1^\top u_1) \Rightarrow u_1^\top \Sigma u_1 = \lambda_1$$

- $\Sigma u_1 = \lambda_1 u_1$  indicates  $u_1$  is an eigenvector of  $\Sigma$  and this can be solved using eigendecomposition.



# Dimensionality Reduction

- In statistics and machine learning, dimensionality reduction is to reduce the number of random variables by obtaining a set of principal variables.
  - ❑ Principal Component Analysis (PCA)
  - ❑ Canonical Correlation Analysis (CCA)
  - ❑ Nonnegative Matrix Factorization (NMF)
  - ❑ Autoencoder
  - ❑ Learning to Hash
  - ❑ Random Projection
  - ❑ Locality-Sensitive Hashing (LSH)
  - ❑ etc.

# Model Selection

- A common problem in machine learning is to select a hyper-parameter which usually determines the structure (or complexity) of the model
  - ❑ Number of components in a GMM
  - ❑ Number of latent dimensions in matrix factorization
  - ❑ Hyper-parameters in neural networks
  - ❑ Hyper-parameters in kernel methods
  - ❑ Hyper-parameters in Bayesian methods
  - ❑ etc.
- Criteria for model selection
  - ❑ Cross-validation – **most frequently used**
  - ❑ Information theory based criteria (AIC, BIC, MDL, etc.)
  - ❑ Bayesian nonparametric methods

# Bayesian Methods

- Recall that Maximum *a posteriori* (MAP) is an estimate of an unknown quantity, that equals the mode of the posterior distribution – **point estimation**.

$$\arg \max_{\theta \in \Theta} p(\theta|D) = \arg \max_{\theta \in \Theta} \frac{p(D|\theta)p(\theta)}{\int_{\theta'} p(D|\theta')p(\theta')d\theta'} = \arg \max_{\theta \in \Theta} p(D|\theta)p(\theta)$$

- Bayesian methods treat parameters as random variables and infer the **posterior distribution**


$$p(\theta|D) = \frac{p(D|\theta)p(\theta)}{\int_{\theta'} p(D|\theta')p(\theta')d\theta'}$$

- Bayesian methods use **predictive distribution** for prediction

$$p(x|D) = \int_{\theta} p(x|\theta)p(\theta|D)d\theta$$

# Point Estimation vs Bayesian Inference

- Point estimation (MLE & MAP) is simpler and more efficient to solve
  - ❑ Build MLE/MAP (objective function) conditioned on  $\theta$
  - ❑ Find optimal  $\theta^*$  using some optimization techniques
  - ❑ Substitute  $\theta^*$  into the model for prediction
- Bayesian inference: There are integrals in both inference stage and prediction stage
- Advantages of Bayesian inference
  - ❑ Avoid local optima
  - ❑ Predict with confidence
  - ❑ Incorporate rich prior knowledge



---

# Thanks

Email: [libin@fudan.edu.cn](mailto:libin@fudan.edu.cn)

# MAP & Generalized Linear Model

- MAP:  $\min_{\theta \in \Theta} \frac{1}{N} \sum_{n=1}^N -\log p(x_n, y_n | \theta) - \log p(\theta)$
- In statistics, the generalized linear model (GLM) is a flexible generalization of ordinary linear regression that allows for response variables that have error distribution models other than a normal distribution.
- GLM generalizes linear regression by allowing the linear model to be related to the target via a **link function**.
  - $p(x_n, y_n | \theta)$  is Gaussian distribution  $\rightarrow$  Link function:  $w^\top x_n = \mu$   
(Identity)  $\rightarrow$  Mean function:  $\mu = w^\top x_n$
  - $p(x_n, y_n | \theta)$  is Logistic function  $\rightarrow$  Link function:  $w^\top x_n = \ln \frac{\mu}{1-\mu}$   
(Logit)  $\rightarrow$  Mean function:  $\mu = \frac{1}{1 + \exp(-w^\top x_n)}$